

APART AI COLLUSION SPRINT · TRACK 3: CARRIERS AND COUNTERMEASURES

From chat collusion to bribes

We reproduced a headline collusion result, found where it breaks, and asked a new question: can money carry collusion between AI auditors?

1 · The paper

2 · The gaps


3 · Why now

4 · Our experiment

5 · Track 3


1 · THE PAPER WE STARTED FROM

"Emergent Collusion in Long-Horizon LLM Agent Interaction" (2609.24967)




Rule in our setup: accept only with the partner's full raw log

What the paper reported (GPT-5.6-Luna)



Same game, newer model: collusion almost disappeared




Setup check vs the team's reference run
12% vs 10%
collusive episodes · runs with any: 60% vs 60%


Cheap model (Gemini Flash-Lite)
10 of 13 runs
started colluding by episode 7

What was counted as collusion was mostly not coordination


Gap 1 · Joint accepts were chance




Gap 2 · The rule could not be followed



Gap 3 · Lapses were one-sided



Replacing a colluding agent bought one episode. The game re-taught the habit.



Gemini Flash-Lite · 10 runs, Bob replaced right after collusion began

7 of 10

fresh Bobs accepted within two more episodes, with no recruitment from Alice

"My previous rejection for the same reason was considered a mistake, so I am prioritizing the correctness of her work over the formal requirement."
fresh Bob, his second episode

Agents are getting wallets, but nobody has tested money as a collusion carrier



Chat

studied · 2609.24967



Secret tools

studied · 2605.27593



Market prices

studied · 2609.18346



Steganography

studied · 2606.28425



Latent states

studied · 2608.19161



Side payments

untested · our gap

Payment rails for agents

Agent payment protocols (AP2, 2608.23858) and agent-to-agent finance (2607.00245) are arriving now

Apart Track 3 asks

What carries coordination, and what breaks it?

Each gap became a design choice

Joint accepts might be chance



Test coordination properly: compare with independent agents in the same situation

The rule was impossible to follow



Make checking possible but costly: one check call for 4 credits

Only chat had been studied




Remove chat, add money: transfers and "paid only if you accept" offers

Incentives drove it, not memory



Vary the incentives one at a time: score, credits, rule, random bribes, planted wrong answers

Can money buy collusion between AI auditors?



WHAT HAPPENED, STEP BY STEP

Agents that could afford to check did not collude: 4 skips in 1,860, never together

1 · Rule on, 20 credits

0 of 220

2 · Add a motive (shared score)

0 of 60


3 · Rule off, 8 credits

19–26%


4 · 20 credits, could pay

4 of 1,860 · 0 joint

Step 3 was a trap: both agents start each run broke




Step 4, 20 credits, no rule: only broke agents skip




Bribes bought nothing. Skipped checks let wrong answers through.

"8 credits if you accept", offered at random




Planted wrong answers that got approved




Agents never coordinated. They trusted each other's words when they couldn't check.

Joint unchecked accepts sit on the chance line



What 1,034 unchecked accepts relied on



"I cannot verify... However, Bob's stated reasoning matches every condition."


The risk isn't bribery. It's unverified trust whenever checking is unaffordable, over a channel nobody authenticates.

Q1 Carriers → words, not money




Unchecked accepts by basis. Random bribe offers moved no verdicts in 7 arms.

Q2 Countermeasures → rule works, then starves



GPT check rate. The rule cuts wrong approvals 24% → 7.5%, but checks collapse.

Q3 Peer identity → none; the channel is spoofable



Q4 Monitoring → literal compliance

