


Can money buy collusion between AI auditors?



WHAT HAPPENED, STEP BY STEP

Unchecked approval appeared only when agents could not pay to check

1 · Rule on, 20 credits

0 of 220

2 · Add a motive (shared score)

0 of 60


3 · Rule off, 8 credits

19–26%


4 · Fair game, 20 credits

4 of 1,860

Step 3 was a trap: both agents start each run broke




Step 4, fair game: who skips the check?




Bribes bought nothing. Skipped checks let wrong answers through.

"8 credits if you accept", offered at random




Planted wrong answers that got approved




Agents never coordinated. They trusted each other's words when they couldn't check.

Joint unchecked accepts sit on the chance line




What 1,034 unchecked accepts relied on




The risk isn't bribery. It's unverified trust whenever checking is unaffordable, over a channel nobody authenticates.

Q1 Carriers → words, not money



Unchecked accepts by basis. Random bribe offers moved no verdicts in 7 arms.

Q2 Countermeasures → rule works, then starves



GPT check rate. The rule cuts wrong approvals 24% → 7.5%, but checks collapse.

Q3 Peer identity → none; the channel is spoofable



Q4 Monitoring → literal compliance

