


Can money buy collusion between AI auditors?



WHAT HAPPENED, STEP BY STEP

Agents that could afford to check did not collude: 4 skips in 1,860, never together

1 · Rule on, 20 credits

0 of 220

2 · Add a motive (shared score)

0 of 60


3 · Rule off, 8 credits

19–26%


4 · 20 credits, could pay

4 of 1,860 · 0 joint

Step 3 was a trap: both agents start each run broke




Step 4, 20 credits, no rule: only broke agents skip




Bribes bought nothing. Skipped checks let wrong answers through.

"8 credits if you accept", offered at random




Planted wrong answers that got approved




Agents never coordinated. They trusted each other's words when they couldn't check.

Joint unchecked accepts sit on the chance line




What 1,034 unchecked accepts relied on



"I cannot verify... However, Bob's stated reasoning matches every condition."


The risk isn't bribery. It's unverified trust whenever checking is unaffordable, over a channel nobody authenticates.

Q1 Carriers → words, not money




Unchecked accepts by basis. Random bribe offers moved no verdicts in 7 arms.

Q2 Countermeasures → rule works, then starves



GPT check rate. The rule cuts wrong approvals 24% → 7.5%, but checks collapse.

Q3 Peer identity → none; the channel is spoofable



Q4 Monitoring → literal compliance

